

Ahmed Awadallah

## Using L1 and L2 Regularization with Superconductivity Dataset

### Introduction

Superconductivity is a state transition where conducting materials go from transmitting current with some resistance(dependent on material) to 0. This occurs at a critical temperature threshold similar to how water will freeze at 0 degrees celsius. Metallic crystals become superconducting generally at low temperatures and it is an active area of research to find more metallic crystals that could be superconductive as well as discover ways to create a superconducting material stable at standard earth temperatures.

### Methods

Our data comes from the DS 7333 Case Study 1. Our objective was to use a linear model with regularization in order to predict the critical temperature as accurately as possible as well as provide the features that proved most helpful to the model.

Personally I added a small twist in the spirit of learning where I produce two models. One that utilized Lasso regularization for feature selection to first shrink the dataset then it was put into a model with Ridge regularization to produce the desired result. The second model simply fed the entire dataset into the model with Ridge regularization.


All models use 5 fold cross validation and our optimizing metric is the root mean squared error.

Dataset contains 168 features, 1 target feature(critical temperature), and 21263 observations. Dataset contains no NAs and contains no glaring issues. The "material" feature was removed as it contained data that was already one hot encoded by the rest of the dataset.


### Results

Overall the rmse of both models are not terrible but not great, upon inspection of the graphs and that we have 21263 observations its likely our model is biased and thus is

incapable of improving any further. Possible future courses of action would likely include exploring if log and polynomial feature transformations may help performance.



Feature Selected Ypred to Ytrue comparison



The all data model had an rmse of 19.15, used an alpha of 1232.85, and took around 22.9 seconds to find the best model.

```
---All Data---  
Best ridge rmse=19.1543, alpha=1232.8467394420684, time=22.9383s  
Top 5 Features and Coefficients  
Ba: 8.5059  
wtd_mean_ThermalConductivity: 5.6796  
wtd_std_Valence: -5.1740  
wtd_std_ThermalConductivity: 4.8497  
Bi: 4.6012
```

The feature selected model had an rmse of 20.38, used an alpha of 932.6, and took around 4.47 seconds to find the best model.

```
---Feature Selected Data---  
Best ridge rmse=20.3756, alpha=932.6033468832219, time=4.4681s  
wtd_std_ThermalConductivity: 9.4389  
Ba: 8.2699  
mean_ElectronAffinity: -4.9017  
wtd_entropy_atomic_mass: 4.8534  
range_atomic_radius: 4.8400
```

The feature selected model performs worse than the full data model. This is to be expected since feature selection removes data from the model that are deemed inconsequential but they still help slightly. In situations where you would want to have optimal performance it may be beneficial to consider leaving everything in

The time it took to train the feature selected was around 5x faster than the full data model as it has less data to train with. In situations where compute power is a limiting factor it may be beneficial to consider taking a small hit to performance in order to reduce the computational load.

Finally it's interesting to note that both models achieved optimal performance at different values of alpha. The all data model opted for a higher alpha than the feature selected model. A higher alpha translates to more regularization so the all data model performed better with stronger penalties to the weights than the model with feature selection in place. This is likely due to the fact that part of the regularization work was carried out by L1 regularization through feature selection so the L2 regularization did not need to be punished as much.

Both models did not give the exact same top 5 important features. Out of the 5 only 2 were shared so this leaves 8 unique features that were deemed important by both models.

- Barium
  - Positive Relationship
  - Present in both models
- wtd\_mean\_ThermalConductivity
  - Positive Relationship
- wtd\_std\_Valence
  - Negative Relationship
- wtd\_std\_ThermalConductivity
  - Positive Relationship
  - Present in both models
- Bismuth
  - Positive Relationship
- mean\_ElectronAffinity
  - Negative Relationship
- Wtd\_entropy\_atomic\_mass
  - Positive Relationship
- range\_atomic\_radius
  - Positive Relationship

All of these features have a strong relationship in at least one of the models so they are worth looking into for future studies. With that said, Barium and wtd\_std\_ThermalConductivity were present in both models which grants them slightly higher priority for investigation than the rest.

## Code

```
1  import pandas as pd
2  import numpy as np
3  import matplotlib.pyplot as plt
4  import sys
5  import time
6  from sklearn.preprocessing import StandardScaler
7  from sklearn.linear_model import Lasso
8  from sklearn.linear_model import Ridge
9  from sklearn.model_selection import cross_val_score
10
11  # Your case study is to build a linear regression model using L1 or L2
12  # regularization (or both) the task to predict the Critical Temperature as
13  # closely as possible. In addition, include in your write-up which variable
14  # carries the most importance.
15
16  INVALID_INPUT = 1
17
18
19  def try_to_read_input_csvs(inputs):
20      try:
21          read_input_csvs(inputs)
22      except Exception:
23          print("This script expects the two csv files associated with Case "
24                "Study1. You either input nothing, 1 file path, or one or more "
25                "of your inputs contains an invalid path.")
26          sys.exit(INVALID_INPUT)
27
28
29  def read_input_csvs(inputs):
30      global df1, df2
31      df1_path = inputs[1]
32      df2_path = inputs[2]
33      df1 = pd.read_csv(df1_path)
34      df2 = pd.read_csv(df2_path)
```

```

37 def linear_cv_feature_selection(x, y, features, print_progress=False):
38     best_alpha_lasso = 0
39     best_rmse_lasso = 0
40     first = True
41     report_freq = 1
42     iter_count = 5
43     param_alpha = 0.01
44     for i in range(0, iter_count):
45         param_alpha = param_alpha*10
46
47         lasso_model = Lasso(alpha=param_alpha)
48
49         rmse_lasso_raw = cross_val_score(lasso_model, x, y, cv=5,
50                                         scoring='neg_root_mean_squared_error')
51
52         rmse_lasso = -1*sum(rmse_lasso_raw)/len(rmse_lasso_raw)
53
54         if first:
55             first = False
56             best_alpha_lasso = param_alpha
57             best_rmse_lasso = rmse_lasso
58         else:
59             if (rmse_lasso < best_rmse_lasso):
60                 best_alpha_lasso = param_alpha
61                 best_rmse_lasso = rmse_lasso
62
63         if ((i+1) % report_freq == 0) and print_progress:
64             print(f"Iteration {i+1}. "
65                 f"Percent done = {(i+1)/iter_count*100:.4f}%")
66
67         # print(f"Best lasso rmse={best_rmse_lasso:.4f} alpha={best_alpha_lasso}")
68         lasso_model = Lasso(alpha=best_alpha_lasso)
69         lasso_model.fit(x, y)
70         nonzero_features = list(lasso_model.coef_ > 0)
71         selected_features = []
72
73         for i in range(len(nonzero_features)):
74             if nonzero_features[i]:
75                 selected_features.append(features[i])
76     return selected_features
77

```

```

29  def linear_cv_compare(x, y, print_progress=False):
30      best_alpha_ridge = 0
31      best_rmse_ridge = 0
32      first = True
33      report_freq = 10
34      iter_count = 100
35      i = 0
36      t1 = time.time()
37      for param_alpha in np.logspace(-6, 6, iter_count):
38          i += 1
39
40          ridge_model = Ridge(alpha=param_alpha)
41
42          rmse_ridge_raw = cross_val_score(ridge_model, x, y, cv=5,
43                                          scoring='neg_root_mean_squared_error')
44
45          rmse_ridge = -1*sum(rmse_ridge_raw)/len(rmse_ridge_raw)
46
47          if first:
48              first = False
49              best_alpha_ridge = param_alpha
50              best_rmse_ridge = rmse_ridge
51          else:
52              if (rmse_ridge < best_rmse_ridge):
53                  best_alpha_ridge = param_alpha
54                  best_rmse_ridge = rmse_ridge
55
56              if ((i+1) % report_freq == 0) and print_progress:
57                  print(f"Iteration {i+1}. "
58                      f"Percent done = {(i+1)/iter_count*100:.4f}%")
59
60      t2 = time.time()
61      elapsed_time = t2-t1
62
63      print(f"Best ridge rmse={best_rmse_ridge:.4f}, "
64          f"alpha={best_alpha_ridge}, time={elapsed_time:.4f}s")
65
66      return best_alpha_ridge

```



```

118 ✓ def get_top_n_coef(feature_list, coef_list, n=5):
119     feature_scores_abs = {}
120     feature_scores = {}
121
122     for i in range(len(feature_list)):
123         feature = feature_list[i]
124         coef = coef_list[i]
125         feature_scores_abs[feature] = abs(coef)
126         feature_scores[feature] = coef
127
128     feature_scores_abs = dict(sorted(feature_scores_abs.items(),
129                                     key=lambda item: item[1], reverse=True))
130     top_n_features = list(feature_scores_abs.keys())[0:n]
131
132     top_feature_scores = {}
133     for feature in top_n_features:
134         top_feature_scores[feature] = feature_scores[feature]
135
136     return top_feature_scores
137
138
139 ✓ if __name__ == '__main__':
140     inputs = sys.argv
141
142     try_to_read_input_csvs(inputs)
143
144     # Combine dataframes
145     y = df1["critical_temp"]
146     df1 = df1.drop(["critical_temp"], axis=1)
147     df2 = df2.drop(["critical_temp", "material"], axis=1)
148
149     # Find which dataframe contains the cts variables that need to be scaled
150     df1_features = list(df1.columns)
151     df2_features = list(df2.columns)
152     if "number_of_elements" in df1_features:
153         cts_features = df1_features
154     else:
155         cts_features = df2_features
156
157     # Scale cts data
158     scaler = StandardScaler()
159     train = pd.concat([df1, df2], axis=1)
160     train[cts_features] = scaler.fit_transform(train[cts_features])

```

```

162 # Lasso Feature Selection
163 print("Lasso Feature Selection Starting. "
164       "This step takes a little bit of time...")
165 selected_features_main = linear_cv_feature_selection(train,
166                                                     y,
167                                                     list(train.columns))
168 train_fs = train[selected_features_main]
169 print("33% Done")
170
171 # Ridge Results
172 print("---All Data---")
173 best_alpha_all_data = linear_cv_compare(train, y)
174 all_data_model = Ridge(alpha=best_alpha_all_data)
175 all_data_model.fit(train, y)
176 all_data_top5 = get_top_n_coef(train.columns, all_data_model.coef_)
177 print("Top 5 Features and Coefficients")
178 for feature in all_data_top5.keys():
179     coef = all_data_top5[feature]
180     print(f"{feature}: {coef:.4f}")
181
182 print("66% Done")
183
184 print("---Feature Selected Data---")
185 best_alpha_fs = linear_cv_compare(train_fs, y)
186 feature_selected_model = Ridge(alpha=best_alpha_fs)
187 feature_selected_model.fit(train_fs, y)
188 feature_selected_top5 = get_top_n_coef(train_fs.columns,
189                                       feature_selected_model.coef_)
190 for feature in feature_selected_top5.keys():
191     coef = feature_selected_top5[feature]
192     print(f"{feature}: {coef:.4f}")
193
194 print("100% Done")
195

```

```
196     # Y_pred to True plots
197     figure, axis = plt.subplots(1, 2, figsize=(16, 10))
198
199     y_pred_all_data = all_data_model.predict(train)
200     axis[0].scatter(y_pred_all_data, y, label="Y Pred v Y True")
201     axis[0].scatter(y, y, label="Example of desired result")
202     axis[0].set_title("All Data Ypred to Ytrue comparison")
203     axis[0].set_xlabel("Y Predicted")
204     axis[0].set_ylabel("Y True")
205     axis[0].legend(loc='upper center', shadow=True)
206
207     y_pred_fs = feature_selected_model.predict(train_fs)
208     axis[1].scatter(y_pred_fs, y, label="Y Pred v Y True")
209     axis[1].scatter(y, y, label="Example of desired result")
210     axis[1].set_title("Feature Selected Ypred to Ytrue comparison")
211     axis[1].set_xlabel("Y Predicted")
212     axis[1].set_ylabel("Y True")
213     axis[1].legend(loc='upper center', shadow=True)
214
215     plt.show()
216
217     print("Program Finished")
218
```